

The Parallel Patterns Library in Visual Studio 2010

Stephan T. Lavavej
Visual C++ Libraries Developer
stl@microsoft.com

Concurrency Is Hard (Let's Go Shopping)

- ❑ Problems
 - Correctness is hard
 - Bugs are notoriously difficult to avoid, find, and fix
 - Performance is hard
 - Efficiency on your hardware
 - Scalability to bigger hardware
 - Elegance is hard
 - When you have low-level primitives and nothing else
 - There is no escape
 - The future is more cores, not (many) more hertz
- ❑ Solution: Do the hard parts in libraries

Parallel Computing Platform

- ❑ Four major parts
 - Concurrency Runtime (ConcRT, "concert")
 - Parallel Patterns Library (PPL, "people")
 - Asynchronous Agents Library
 - Parallel Debugging and Profiling
- ❑ Implemented within `msvcr100.dll/msvcpr100.dll` and `libcmt.lib/libcpmt.lib`
 - Nothing more to link against or redistribute
 - x86 DLLs are currently $735 + 416 = 1151$ KB
 - In VS 2008 SP1 (VC9 SP1), $641 + 560 = 1201$ KB
 - PPL and Agents are mostly header-only

Concurrency Runtime (ConcRT)

- ❑ What happens when a single process contains independent concurrent components?
- ❑ ConcRT...
 - Arbitrates between multiple requests for computing resources within a single process
 - Can reclaim resources when its cooperative blocking mechanisms are used
 - Is aware of locality (e.g. cores sharing caches)
 - Is aware of Windows 7 User Mode Scheduled Threads
 - Isn't used by application developers directly
 - Is used by library developers to build programming models
 - PPL, Agents, Boost 1.42.0?

Parallel Patterns Library (PPL)

- ❑ Namespace
 - Concurrency
- ❑ Tasks
 - task_handle
 - task_group
 - structured_task_group
- ❑ Algorithms
 - parallel_invoke()
 - parallel_for()
 - parallel_for_each()
- ❑ Containers
 - combinable
 - concurrent_queue
 - concurrent_vector
- ❑ Synchronization Primitives
 - critical_section
 - event
 - reader_writer_lock

Tasks

- ❑ Task (`task_handle`)
 - Stores a given functor
 - Represents a sub-computation that can be executed concurrently with other tasks
- ❑ Task Group (`task_group`)
 - Stores a given bunch of tasks
 - Performs a computation by executing its tasks concurrently and waiting for them to finish
- ❑ A task can use a task group to execute nested tasks
- ❑ Structured Task Group (`structured_task_group`)
 - Less overhead, more restrictions

Tasks Example: Serial Recursion

```
void quicksort(vector<int>::iterator first,
vector<int>::iterator last) {
    if (last - first < 2) { return; }
    int pivot = *first;
    auto mid1 = partition(first, last,
        [=](int elem) { return elem < pivot; });
    auto mid2 = partition( mid1, last,
        [=](int elem) { return elem == pivot; });
    quicksort(first, mid1);
    quicksort(mid2, last);
}
```

Tasks Example: Parallel Recursion

```
void quicksort(vector<int>::iterator first,
vector<int>::iterator last) {
    if (last - first < 2) { return; }
    int pivot = *first;
    auto mid1 = partition(first, last,
        [=](int elem) { return elem < pivot; });
    auto mid2 = partition( mid1, last,
        [=](int elem) { return elem == pivot; });
    task_group g;
    g.run([=] { quicksort(first, mid1); });
    g.run([=] { quicksort(mid2, last); });
    g.wait();
}
```

Tasks Example: Performance

- ❑ Intel Core 2 Quad Q9450 (Yorkfield 2.66 GHz)
- ❑ 50,000,000 elements
- ❑ Serial: 2939.85 ms
- ❑ Parallel: 1308.12 ms
- ❑ Speedup: 2.247

Algorithms: parallel_invoke()

- ❑ Takes 2 to 10 functors
- ❑ Executes them concurrently
- ❑ Waits for them to finish
- ❑ Example:

```
parallel_invoke(  
    [=] { quicksort(first, mid1); },  
    [=] { quicksort(mid2, last); }  
);
```

parallel_invoke() Performance

- ❑ Intel Core 2 Quad Q9450 (Yorkfield 2.66 GHz)
- ❑ 50,000,000 elements
- ❑ Serial: 2939.85 ms
- ❑ task_group
 - 1308.12 ms
 - Speedup: 2.247
- ❑ parallel_invoke()
 - 1122.9 ms
 - Speedup: 2.618

Algorithms: parallel_for()

- ❑ Usage:
 - `parallel_for(first, last, functor);`
 - `parallel_for(first, last, step, functor);`
 - Requires `step > 0`
- ❑ Concurrently calls functor with each index in
[`first`, `last`)

Containers: combinable

- ❑ Collects results from tasks
- ❑ Merges them into a final result
- ❑ A lock-free alternative to a shared variable
- ❑ Notable `combinable<T>` members:
 - `combinable()`
 - `combinable(generator)`
 - `T& local()`
 - `T combine(combiner)`
 - `combine_each(accumulator)`

parallel_for() / combinable

Example: Serial Iteration

```
vector<int> v;

for (int i = 2; i < 5000000; ++i) {
    if (is_carmichael(i)) {
        v.push_back(i);
    }
}
```

parallel_for() / combinable

Example: Parallel Iteration

```
vector<int> v;  
  
combinable<vector<int>> c;  
  
parallel_for(2, 5000000, [&](int i) {  
    if (is_carmichael(i)) {  
        c.local().push_back(i);  
    }  
});  
  
c.combine_each([&](const vector<int>& sub) {  
    v.insert(v.end(), sub.begin(), sub.end());  
});  
  
sort(v.begin(), v.end());
```

parallel_for() / combinable

Example: Performance

- ❑ Intel Core 2 Quad Q9450 (Yorkfield 2.66 GHz)
- ❑ 4,999,998 indices
- ❑ Serial: 8679.61 ms
- ❑ Parallel: 2183.43 ms
- ❑ Speedup: 3.975

Containers: concurrent_queue And concurrent_vector

- ❑ concurrent_queue and concurrent_vector are...
 - Lock-free data structures
 - Aside: shared_ptr is also lock-free
 - Similar to queue and vector
 - NOT IDENTICAL
 - Example: vector is contiguous, concurrent_vector isn't
 - Suspiciously familiar
 - Intel Threading Building Blocks
 - Not in VS 2010 Beta 1 (VC10 Beta 1)
- ❑ concurrent_vector<T> can be better than combinable<vector<T>>
 - combinable<int>, etc. is still useful

concurrent_vector

Example: Code And Performance

```
concurrent_vector<int> c;

parallel_for(2, 5000000, [&](int i) {
    if (is_carmichael(i)) {
        c.push_back(i);
    }
});

sort(c.begin(), c.end());
```

- ❑ combinable: 2183.43 ms
- ❑ concurrent_vector: 2181.02 ms

Algorithms: parallel_for_each()

- Usage:
 - `parallel_for_each(first, last, functor);`
- Concurrently calls functor with each element in `[first, last)`
- Accepts forward iterators (e.g. `forward_list`)
- Really likes random access iterators (e.g. `vector`)

`parallel_for_each()`

Example: Performance

- ❑ Intel Core 2 Quad Q9450 (Yorkfield 2.66 GHz)
- ❑ 4,999,998 elements
- ❑ `forward_list`
 - Serial: 8681.78 ms
 - Parallel: 2680.05 ms
 - Speedup: 3.239
- ❑ `vector`
 - Serial: 8682.49 ms
 - Parallel: 2189.59 ms
 - Speedup: 3.965

Synchronization Primitives

- ❑ Provided by <concr.h>, not <ppl.h>:
 - critical_section
 - event
 - reader_writer_lock
- ❑ These cooperative blocking mechanisms talk to ConcRT and allow it to schedule something else
- ❑ Windows API synchronization primitives still work
 - But ConcRT isn't aware of them, so they're less efficient
 - Except on Windows 7, thanks to UMS Threads
 - Still, you should prefer ConcRT's synchronization primitives

Scaling: Quad-Core Vs. Octo-Core Speedups

- Intel Core 2 Quad Q9450 (Yorkfield 2.66 GHz)
- Intel Xeon E5335 (8 cores, Clovertown 2.00 GHz)

Example	Quad	Octo
task_group	2.247	3.455
parallel_invoke()	2.618	3.857
combinable	3.975	7.991
concurrent_vector	3.980	7.988
p_f_e() forward_list	3.239	6.242
p_f_e() vector	3.965	7.965

Questions?

- For more information, see:
 - msdn.com/concurrency
 - channel9.msdn.com/tags/Parallelism
 - blogs.msdn.com/vcblog
- stl@microsoft.com

Bonus Slides!

Tasks Example (1/5)

```
C:\Temp>type quicksort.cpp
#include <algorithm>
#include <iostream>
#include <ostream>
#include <vector>
#include <ppl.h>
#include <windows.h>
using namespace std;
using namespace Concurrency;

long long counter() {
    LARGE_INTEGER li;
    QueryPerformanceCounter(&li);
    return li.QuadPart;
}

long long frequency() {
    LARGE_INTEGER li;
    QueryPerformanceFrequency(&li);
    return li.QuadPart;
}
```

Tasks Example (2/5)

```
void quicksort(vector<int>::iterator first, vector<int>::iterator last) {
    if (last - first < 2) {
        return;
    }

    int pivot = *first;

    auto mid1 = partition(first, last, [=](int elem) { return elem < pivot; });
    auto mid2 = partition( mid1, last, [=](int elem) { return elem == pivot; });

#ifndef USE_PPL
    task_group g;

    g.run([=] { quicksort(first, mid1); });
    g.run([=] { quicksort(mid2, last); });

    g.wait();
#else
    quicksort(first, mid1);
    quicksort(mid2, last);
#endif
}
```

Tasks Example (3/5)

```
int main() {
    vector<int> v;

    for (int k = 1, n = 1; v.size() < 50000000; ) {
        v.push_back(n);

        if (n == 1 || n > 70000000) {
            n = ++k;
        } else if (n % 2 == 0) {
            n /= 2;
        } else {
            n = 3 * n + 1;
        }
    }
}
```

Tasks Example (4/5)

```
long long start = counter();

quicksort(v.begin(), v.end());

long long finish = counter();

cout << (finish - start) * 1000.0 / frequency()
    << " ms" << endl;

if (is_sorted(v.begin(), v.end())) {
    cout << "SUCCESS" << endl;
} else {
    cout << "FAILURE" << endl;
}

}
```

Tasks Example (5/5)

```
C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0  
    /Fequicksort_serial.exe quicksort.cpp
```

quicksort.cpp

Generating code

Finished generating code

```
C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0  
    /DUSE_PPL /Fequicksort_parallel.exe quicksort.cpp
```

quicksort.cpp

Generating code

Finished generating code

```
C:\Temp>quicksort_serial
```

2939.85 ms

SUCCESS

```
C:\Temp>quicksort_parallel
```

1308.12 ms

SUCCESS

parallel_for() / combinable

Example (1/5)

```
C:\Temp>type carmichael.cpp
#include <algorithm>
#include <iostream>
#include <ostream>
#include <vector>
#include <ppl.h>
#include <windows.h>
using namespace std;
using namespace Concurrency;

long long counter() {
    LARGE_INTEGER li;
    QueryPerformanceCounter(&li);
    return li.QuadPart;
}

long long frequency() {
    LARGE_INTEGER li;
    QueryPerformanceFrequency(&li);
    return li.QuadPart;
}
```

parallel_for() / combinable

Example (2/5)

```
bool is_carmichael(const int n) {
    if (n < 2) { return false; }

    int k = n;

    for (int i = 2; i <= k / i; ++i) {
        if (k % i == 0) {
            if ((k / i) % i == 0) { return false; }
            if ((n - 1) % (i - 1) != 0) { return false; }
            k /= i;
            i = 1;
        }
    }

    return k != n && (n - 1) % (k - 1) == 0;
}
```

parallel_for() / combinable

Example (3/5)

```
int main() {
    vector<int> v;

    long long start = counter();

#ifndef USE_PPL
    combinable<vector<int>> c;

    parallel_for(2, 5000000, [&](int i) {
        if (is_carmichael(i)) {
            c.local().push_back(i);
        }
    });
    c.combine_each([&](const vector<int>& sub) {
        v.insert(v.end(), sub.begin(), sub.end());
    });
    sort(v.begin(), v.end());
#else
```

parallel_for() / combinable

Example (4/5)

```
for (int i = 2; i < 5000000; ++i) {
    if (is_carmichael(i)) {
        v.push_back(i);
    }
}
#endif

long long finish = counter();

cout << (finish - start) * 1000.0 / frequency() << " ms" << endl;

cout << v.size() << " Carmichael numbers found." << endl;

cout << "First five: ";
for_each(v.begin(), v.begin() + 5, [](int i) { cout << i << " "; });
cout << endl;

cout << "Last five: ";
for_each(v.end() - 5, v.end(), [](int i) { cout << i << " "; });
cout << endl;
}
```

parallel_for() / combinable Example (5/5)

```
C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0  
    /Fecarmichael_serial.exe carmichael.cpp  
carmichael.cpp  
Generating code  
Finished generating code
```

```
C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0 /DUSE_PPL  
    /Fecarmichael_parallel.exe carmichael.cpp  
carmichael.cpp  
Generating code  
Finished generating code
```

```
C:\Temp>carmichael_serial  
8679.61 ms  
74 Carmichael numbers found.  
First five: 561 1105 1729 2465 2821  
Last five: 4335241 4463641 4767841 4903921 4909177
```

```
C:\Temp>carmichael_parallel  
2183.43 ms  
74 Carmichael numbers found.  
First five: 561 1105 1729 2465 2821  
Last five: 4335241 4463641 4767841 4903921 4909177
```

concurrent_vector

Example (1/4)

```
C:\Temp>type convector.cpp
#include <algorithm>
#include <iostream>
#include <ostream>
#include <concurrent_vector.h>
#include <ppl.h>
#include <windows.h>
using namespace std;
using namespace Concurrency;

long long counter() {
    LARGE_INTEGER li;
    QueryPerformanceCounter(&li);
    return li.QuadPart;
}

long long frequency() {
    LARGE_INTEGER li;
    QueryPerformanceFrequency(&li);
    return li.QuadPart;
}
```

concurrent_vector

Example (2/4)

```
bool is_carmichael(const int n) {
    if (n < 2) { return false; }

    int k = n;

    for (int i = 2; i <= k / i; ++i) {
        if (k % i == 0) {
            if ((k / i) % i == 0) { return false; }
            if ((n - 1) % (i - 1) != 0) { return false; }
            k /= i;
            i = 1;
        }
    }

    return k != n && (n - 1) % (k - 1) == 0;
}
```

concurrent_vector

Example (3/4)

```
int main() {
    long long start = counter();

    concurrent_vector<int> c;

    parallel_for(2, 5000000, [&](int i) {
        if (is_carmichael(i)) {
            c.push_back(i);
        }
    });
    sort(c.begin(), c.end());

    long long finish = counter();

    cout << (finish - start) * 1000.0 / frequency() << " ms" << endl;

    cout << c.size() << " Carmichael numbers found." << endl;
```

concurrent_vector

Example (4/4)

```
cout << "First five: ";
for_each(c.begin(), c.begin() + 5, [](int i) { cout << i << " "; });
cout << endl;

cout << "Last five: ";
for_each(c.end() - 5, c.end(), [](int i) { cout << i << " "; });
cout << endl;
}
```

```
C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0
    convector.cpp
convector.cpp
Generating code
Finished generating code
```

```
C:\Temp>convector
2181.02 ms
74 Carmichael numbers found.
First five: 561 1105 1729 2465 2821
Last five: 4335241 4463641 4767841 4903921 4909177
```

parallel_for_each()

Example (1/5)

```
C:\Temp>type pfe.cpp
#include <algorithm>
#include <forward_list>
#include <iostream>
#include <numeric>
#include <ostream>
#include <vector>
#include <concurrent_vector.h>
#include <ppl.h>
#include <windows.h>
using namespace std;
using namespace Concurrency;

long long counter() {
    LARGE_INTEGER li;
    QueryPerformanceCounter(&li);
    return li.QuadPart;
}

long long frequency() {
    LARGE_INTEGER li;
    QueryPerformanceFrequency(&li);
    return li.QuadPart;
}
```

parallel_for_each()

Example (2/5)

```
bool is_carmichael(const int n) {
    if (n < 2) { return false; }

    int k = n;

    for (int i = 2; i <= k / i; ++i) {
        if (k % i == 0) {
            if ((k / i) % i == 0) { return false; }
            if ((n - 1) % (i - 1) != 0) { return false; }
            k /= i;
            i = 1;
        }
    }

    return k != n && (n - 1) % (k - 1) == 0;
}
```

parallel_for_each()

Example (3/5)

```
int main() {  
  
#ifdef FORWARD  
    forward_list<int> src(4999998);  
#endif  
  
#ifdef RANDOM  
    vector<int> src(4999998);  
#endif  
  
    iota(src.begin(), src.end(), 2);  
  
    long long start = counter();
```

parallel_for_each()

Example (4/5)

```
#ifdef SERIAL
    vector<int> dest;

    for_each(src.begin(), src.end(),
        [&](int i) {
            if (is_carmichael(i)) {
                dest.push_back(i);
            }
        });
#endif

#ifndef PARALLEL
    concurrent_vector<int> dest;

    parallel_for_each(src.begin(), src.end(),
        [&](int i) {
            if (is_carmichael(i)) {
                dest.push_back(i);
            }
        });
}

sort(dest.begin(), dest.end());
#endif
```

parallel_for_each()

Example (5/5)

```
long long finish = counter();

cout << (finish - start) * 1000.0 / frequency() << " ms" << endl;

cout << dest.size() << " Carmichael numbers found." << endl;

cout << "First five: ";
for_each(dest.begin(), dest.begin() + 5, [](int i) { cout << i << " "; });
cout << endl;

cout << "Last five: ";
for_each(dest.end() - 5, dest.end(), [](int i) { cout << i << " "; });
cout << endl;
}

cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0 pfe.cpp /DFORWARD /DSERIAL
    /Fepfe_forward_serial.exe
cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0 pfe.cpp /DFORWARD /DPARALLEL
    /Fepfe_forward_parallel.exe
cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0 pfe.cpp /DRANDOM /DSERIAL
    /Fepfe_random_serial.exe
cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0 pfe.cpp /DRANDOM /DPARALLEL
    /Fepfe_random_parallel.exe
```

Ultrasort Example (1/7)

```
C:\Temp>type ultrasort.cpp
#include <stddef.h>
#include <algorithm>
#include <iostream>
#include <memory>
#include <ostream>
#include <tuple>
#include <vector>
#include <ppl.h>
#include <windows.h>
using namespace std;
using namespace Concurrency;

long long counter() {
    LARGE_INTEGER li;
    QueryPerformanceCounter(&li);
    return li.QuadPart;
}

long long frequency() {
    LARGE_INTEGER li;
    QueryPerformanceFrequency(&li);
    return li.QuadPart;
}
```

Ultrasort Example (2/7)

```
void ultrasort(vector<int>::iterator first, vector<int>::iterator last) {  
    size_t elems = last - first;  
  
    unsigned int procs = CurrentScheduler::Get()  
        ->GetNumberOfVirtualProcessors();  
  
    if (elems < procs) {  
        sort(first, last);  
        return;  
    }  
  
    size_t slice = elems / procs;  
  
    typedef tuple<vector<int>::iterator, vector<int>::iterator,  
        shared_ptr<event>> range_t;  
  
    vector<range_t> ranges;  
  
    task_group tasks;
```

Ultrasort Example (3/7)

```
for (unsigned int i = 0; i < procs; ++i) {
    auto a = first + slice * i;
    auto b = i + 1 < procs
        ? first + slice * (i + 1)
        : last;

    auto e = make_shared<event>();

    ranges.push_back(make_tuple(a, b, e));

    tasks.run([=] {
        sort(a, b);
        e->set();
    });
}
```

Ultrasort Example (4/7)

```
while (ranges.size() > 1) {
    vector<range_t> fused;

    for (size_t i = 0; i + 1 < ranges.size(); i += 2) {
        auto a0 = get<0>(ranges[i]);
        auto b0 = get<1>(ranges[i]);
        auto e0 = get<2>(ranges[i]);

        auto b1 = get<1>(ranges[i + 1]);
        auto e1 = get<2>(ranges[i + 1]);

        auto e = make_shared<event>();

        fused.push_back(make_tuple(a0, b1, e));

        tasks.run([=] {
            event * both[] = { e0.get(), e1.get() };
            event::wait_for_multiple(both, 2, true);
            inplace_merge(a0, b0, b1);
            e->set();
        });
    }
}
```

Ultrasort Example (5/7)

```
if (ranges.size() % 2 != 0) {
    fused.push_back(ranges.back());
}

ranges.swap(fused);
}

tasks.wait();
}

int main() {
    vector<int> v;

    for (int k = 1, n = 1; v.size() < 500000000; ) {
        v.push_back(n);

        if (n == 1 || n > 700000000) {
            n = ++k;
        }
    }
}
```

Ultrasort Example (6/7)

```
    } else if (n % 2 == 0) {
        n /= 2;
    } else {
        n = 3 * n + 1;
    }
}
```

```
long long start = counter();
```

```
#ifdef USE_PPL
    ultrasort(v.begin(), v.end());
#else
    sort(v.begin(), v.end());
#endif
```

```
long long finish = counter();
```

Ultrasort Example (7/7)

```
cout << (finish - start) * 1000.0 / frequency()
    << " ms" << endl;

if (is_sorted(v.begin(), v.end())) {
    cout << "SUCCESS" << endl;
} else {
    cout << "FAILURE" << endl;
}

}

C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0
    /Feultrasort_serial.exe ultrasort.cpp
ultrasort.cpp
Generating code
Finished generating code

C:\Temp>cl /EHsc /nologo /W4 /MT /O2 /GL /D_ITERATOR_DEBUG_LEVEL=0 /DUSE_PPL
    /Feultrasort_parallel.exe ultrasort.cpp
ultrasort.cpp
Generating code
Finished generating code
```

Ultrasort Example: Scaling (Or Lack Thereof)

- ❑ Intel Core 2 Quad Q9450 (Yorkfield 2.66 GHz)
 - parallel_invoke(): **2.618** speedup
 - ultrasort(): **2.416** (3286.28 ms to 1360.12 ms)
- ❑ Intel Xeon E5335 (8 cores, Clovertown 2.00 GHz)
 - parallel_invoke(): **3.857** speedup
 - ultrasort(): **3.192** (4340.89 ms to 1359.94 ms)
- ❑ Find the inefficiency!