# **Boost + Software Transactional Memory**

Boost Library Conference (BoostCon)
May 5, 2009





# Justin E. Gottschlich, Ph.D. Student

Department of Electrical and Computer Engineering
University of Colorado-Boulder

#### **Outline**

#### Purpose

#### **Problem**

- § The Multicore Revolution
- § Mutual Exclusion

#### **Using TM**

§ Five Interfaces + Beyond

**Under the Hood** 

Roadmap to boost::stm







#### Collaborators



Jeremy G. Siek

Manish Vachharajani



Vicente J. Botet Escriba



Maurice Herlihy











# The Multicore Revolution

# **And Our Free Lunch**





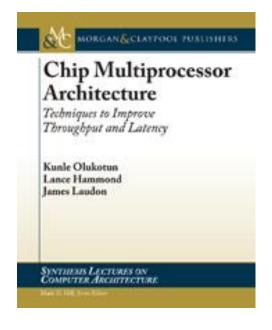


# Why Multicore Computers?

# Manufacturers could not make uniprocessors faster

Heat dissipation a problem

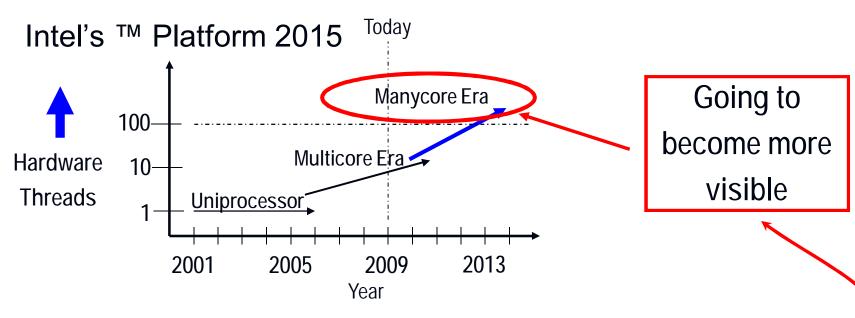
"Large uniprocessors are no longer scaling in performance ..."







#### "The Free (Performance) Lunch Is Over"



#### Disadvantages:

- § No universal performance benefit like unicore
- Only parallel software benefits





# **Parallelizing Code**

#### How do we make software parallel?

- § Rewrite code w/ parallel libraries
  - § boost::thread, boost::mutex, boost::mpi

#### So what **is** the problem?

§ Conventional parallel programming is very difficult





#### **Bank Account Problem**

#### Checking account

- § Global int C
- § Controlled by lock LC

#### Savings account

- § Global int S
- § Controlled by lock LS

#### Goal:

- § Move \$\$ checking to savings
- § Don't bounce checks!
- § Atomic Operation?





# Library Approach: Inconsistent Data

#### library interface boundary

```
void setCandS void getCandS void getSandC
(int v1, int v2) (int &v1, int &v2) (int &v2, int &v1)
{
    setC(v1);
    getC(v1);
    getS(v2);
    getC(v1);
}
```



can lead to inconsistent data



### Hacker's Approach: Deadlock

```
void setCandS
                   void getCandS
                                       void getSandC
                   (int &c, int &s)
(int c, int s)
                                       (int &s, int &c)
  lock(LC);
                                         lock(LS);
                     lock(LC);
                                         lock(LC);
  lock(LS);
                     lock(LS);
  C = C;
                     c = C;
                                         s = S;
  S = s;
                     s = S;
                                         c = C;
  unlock(LC);
                     unlock(LC);
                                         unlock(LS);
  unlock(LS);
                     unlock(LS);
                                         unlock(LC);
```





## **Composing Parallel Libraries**

```
void setC(int c)
                    void getC(int &c)
                                       void setS(int s)
                                                           void getS(int &s)
  atomic(t) {
                      atomic(t) {
                                          atomic(t) {
                                                             atomic(t) {
  t.w(C) = c;
                      c = t.r(C);
                                         t.w(S) = s;
                                                              s = t.r(S);
                      } end_atom
  } end atom
                                          } end atom
                                                             } end_atom
```

#### Real TBoost.STM Code

```
void setCandS
                      void getCandS
                                             void getSandC
(int c, int s)
                      (int &c, int &s)
                                             (int &s, int &c)
  atomic(t) {
                        atomic(t) {
                                               atomic(t) {
    setC(c);
                          getC(c);
                                                 getS(s);
    setS(s);
                          getS(s);
                                                 getC(c);
  } end_atom
                        } end_atom
                                               } end_atom
```





#### **Brief Overview of TM**

#### Transactional Memory (TM)

- § Optimistic concurrency
- § Synchronization: transaction
- § Transactions compose
- § Generally retried until commit

#### Transaction (tx): finite set of operations

- § Requires ACI
  - § Atomic: all or nothing
  - § Consistent: only legal memory states
  - § Isolated: other txes cannot see until committed

#### Contention Management (CM)

§ Mechanism for tx forward progress





## **Example Transaction and ACI**

```
native_trans<int> x = 0, y = 0;
Thread 1
                              Thread 2
                              atomic(t) {
atomic(t) {
  ++t.w(x);
                         Transaction terminated
  --t.w(y);
} end_atem
                                end_atom
   CM: Thread 1's tx commits, Thread 2's tx aborts
   end state, x = 1, y = -1, <u>atomic</u> all or nothing
   abort if x or y change, ensure <u>consist</u>ent
   mid-state, x = 1, y = 0 is isolated
```





# TBoost.STM In Five Interfaces (And Two Class Templates)







### TBoost.STM's native\_trans

#### native\_trans<>

- § Used for all native types
- § Arrays ok, not pointers

#### **Examples:**

```
§ native_trans<bool> txBool;
§ native_trans<float> txFloat;
§ native_trans<int> txIntArr[100];
```





#### **TBoost.STM** in Five Interfaces

```
void transaction::initialize()
§ Initializes STM, always first call
void transaction::initialize_thread()
§ Initializes thread, call per thread before using tx
atomic(<transaction name>)
{ <compound> } end_atom
§ Constructs a transaction object
§ Executes <compound> until successful
template <typename T> T&
transaction::w(T&)
§ Perform tx write; returns ref of object to write
template <typename T> T const&
transaction::r(T const&)
§ Perform tx read; returns const ref of object to read
               <> = user supplied
```

# Using native\_trans<int>

```
native_trans<int> C = 1000;
int deposit_and_balance()
                                 Starting State:
                                 § C = 1000
  int c = 100, bal = -1;
                                 § bal = -1
  atomic (t)
                                 Ending State:
    t.w(C) += C;
                                 § C = 1100
    bal = t.r(C);
                                 § bal = 1100
  } end_ Isolated State Captured
  return bal;
```





# Using native\_trans<int> []

```
native trans<int> arr[100];
                                      Starting State:
                                      s arr = 0, 0, 0, ...
void read arr(int out[])
  atomic(t) {
                                      Ending State:
    for (int i = 0; i < 100; ++i)
      out[i] = t.r(arr[i]);
                                       s arr = 0, 1, 2, ...
   end atom
void write_arr()
                                      Note: read arr()
  atomic(t) {
                                       § Never sees:
    for (int i = 0; i < 100; ++i)
                                          § 0, 1, 2, 0, 0, ...
      t.w(arr[i]) = i;
                                       § Only: 0, 0, 0, ...
   end_atom
                                       § Or: 0, 1, 2, 3, ...
```





### TBoost.STM's transaction\_object

#### transaction\_object

- § Used for all user-defined types
- § Must be default and copy constructable
- § Must have callable operator=
- § Uses curiously recurring template





# Building transaction\_object (char\*)

```
class str_trans : public transaction_object <str_trans>
  char *b;
public:
   str trans() : b (0) {}
   str_trans(str_trans const &rhs) : b_(0) { cp(rhs.b_); }
   str trans& operator=(str trans const &rhs)
   { cp(rhs.b_); return *this; }
   void cp(char const *buf)
     delete[]b;b=0;
      if (0 == buf) return;
     b_ = new char [strlen(buf)+1]; strcpy(b_, buf);
   char const * buf() const { return b_ ? b_ : ""; }
   ~str_trans() { delete [] b_; }
```



# Building transaction\_object (char\*)

```
using namespace std;
str_trans str;
void read_str(string &s)
   atomic(t) {
      s = t.r(str).buf();
   } end_atom
void write_str(string const &s)
   atomic(t) {
     t.w(str).cp(s.c_str());
   } end_atom
```



# TBoost.STM Beyond Five Interfaces







# How Do We Perform I/O (or Any Irreversible Operation)?

```
transaction::make_irrevocable()
§ If fails, throws exception
§ Otherwise, guaranteed to commit
     void hello_world() {
       atomic(t) {
         write str((string)
            "hello concurrent world");
         string s; read str(s);
         t.make irrevocable();
         cout << s.c str();</pre>
       } end atom
```





# **Building lists**

```
template <typename T> class list node :
  public transaction_object< list_node<T> >
  list node *next;
  T value;
};
template <typename T> class linked list
public:
  bool lookup(T const &val) const; //assume coded
  bool insert(list node<T> const &rhs); // new
  bool remove(list node<T> const &rhs); // del
private:
  list node<T> head ;
};
```



#### new and delete

Do not use new and delete for txes!

```
To create memory
```

```
§ T* new memory()
```

```
§ T* new_memory_copy(T const &rhs)
```

#### To destroy memory

```
§ delete_memory(T &rhs)
```





# **Building list::insert**

```
bool insert(list_node<T> const &rhs) {
   atomic(t) {
      list_node<T> const *headP = &t.r(head_);
      if (NULL != headP->next()) {
         // find location to insert
         list node<T> *newN = t.new memory copy(rhs);
         // set prev, new and next
      } else {
         list node<T> *newN = t.new memory copy(rhs);
         // set head
   } end_atom
   return true;
```





## **Building list::remove**

```
bool remove(list_node<T> const &rhs) {
  atomic(t) {
    // find the item to remove
    if (cur->value() == rhs.value()) {
      // reset pointers
      t.delete_memory(*cur);
      t.end(); return true;
   end_atom
                       what does this do?
 return false;
```





### before\_retry

```
atomic(<tx>) { <compound> }
before_retry { <compound> }
```

- § before\_retry executes before tx is retried
- § useful for cleanup
- § changing tx priority, CM strategy, etc.





## **Building algo::move**

```
template <typename C> bool algo::move
(C &l1, C &l2, C::node_type const &v)
 bool r = false;
  atomic(t) {
    if (11.lookup(v) && !12.lookup(v))
      11.remove(v); 12.insert(v); r = true;
  } before_retry { r = false; }
  return r;
```





#### References and Smart Pointers

#### Generally, don't do this:

```
§ native_trans<int> &x = t.w(X);
§ a bad idea
§ native_trans<int> &x = t.r(X);
§ an even worse idea!
```

#### Instead, use:

```
§ boost::stm::read_ptr(transaction &t, T)
§ boost::stm::write_ptr(transaction &t, T)
```





# Problems With Holding References

```
native_trans<int> txInt = 0;
void test_parent() {
  atomic(t) {
    native_trans<int> &tx = t.r(txInt);
    if (0 == tx) test_nested();
    cout << tx << endl;
  } end_atom
void test_nested() {
  atomic(t) { ++t.w(txInt); } end_atom
```





### Use read\_ptr

```
native_trans<int> txInt = 0;
void test_parent() {
  atomic(t) {
    read_ptr< native_trans<int> >
      tx(t, txInt);
    if (0 == *tx) test_nested();
    cout << *tx << endl;
  } end_atom
void test_nested() {
  atomic(t) { ++t.w(txInt); } end_atom
```





### **Use write\_ptr**

```
native_trans<int> txInt = 0;
void test_parent() {
  atomic(t) {
    write_ptr< native_trans<int> >
      tx(t, txInt);
    if (0 == *tx) test nested();
    cout << ++*tx << endl;
  } end_atom
void test_nested() {
  atomic(t) { ++t.w(txInt); } end_atom
```







# **Under the Hood:**

Fast Synrchonization







# **Optimistic Concurrency**

#### Pessimistic concurrency (locks)

§ Critical section limited to single thread

#### Optimistic concurrency (transactions)

- § Unlimited threads can execute critical section
- § Truly optimistic concurrency
  - § Simultaneously conflicting critical section execution
  - § Not supported by all TMs

#### Important aspects

- § Consistency checking
- § Updating policies





# **Consistency Checking**

#### Consistency checking

- § Process to verify state of transaction is consistent
- § Inconsistency:  $C_w \cap (I_w \cup I_R) \neq \emptyset$

#### Validation

- § Each memory location has version #
- § When committing
  - § tx verifies version # is same
  - § increments version for writes

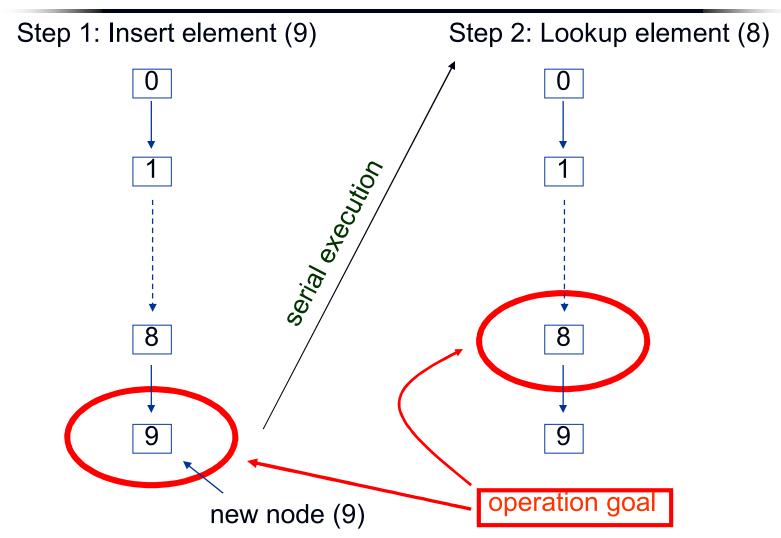
#### Invalidation

- § No version #s used
- § When committing
  - § abort txes reading/writing mem its writing





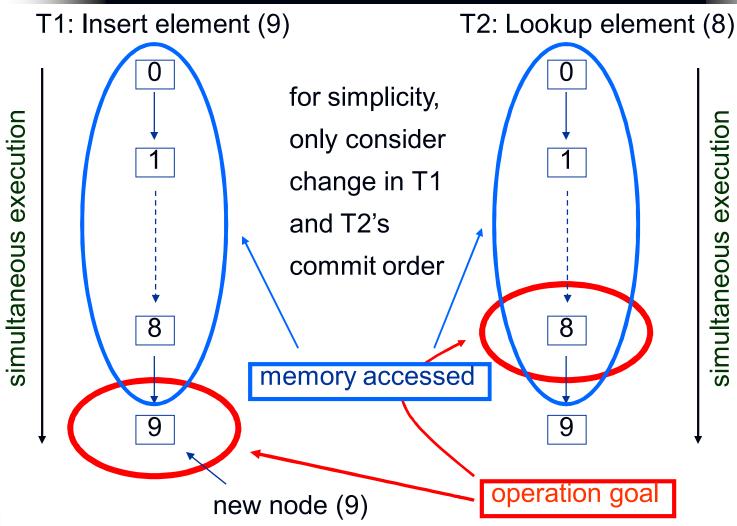
# Linked List (Pessimistic)







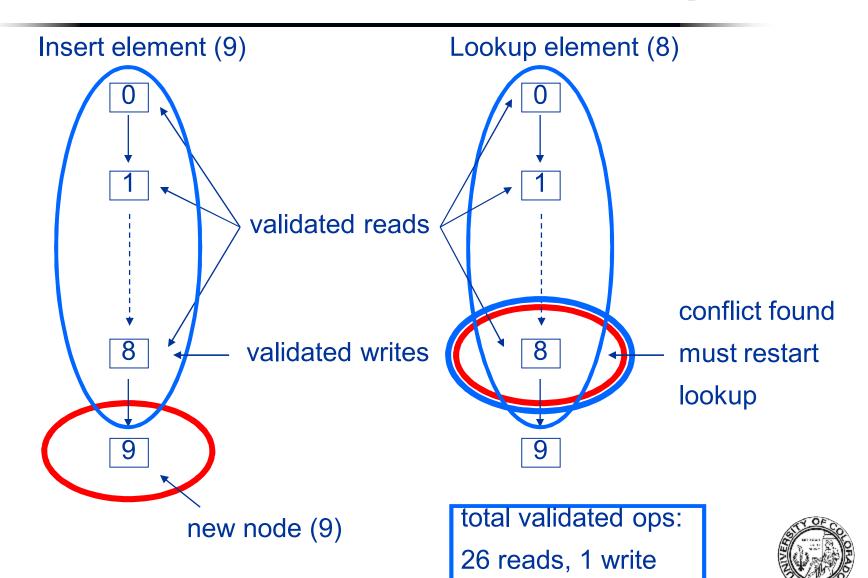
# Linked List (Optimistic)





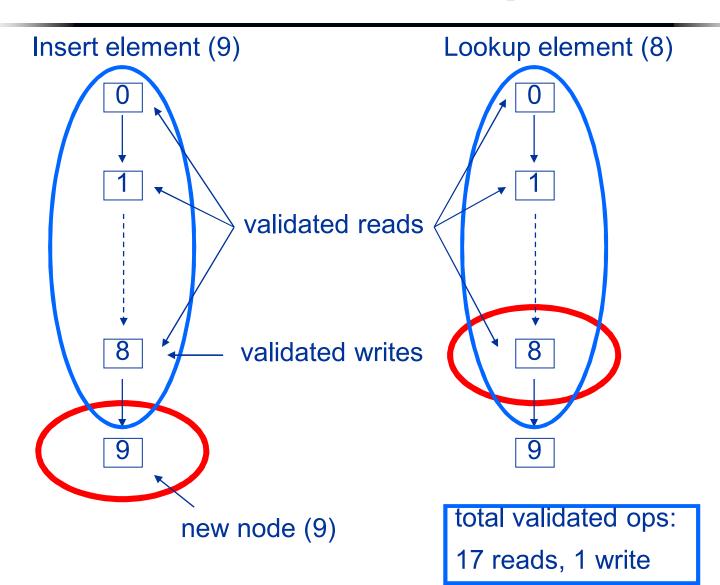


## Validation – Insert, Lookup





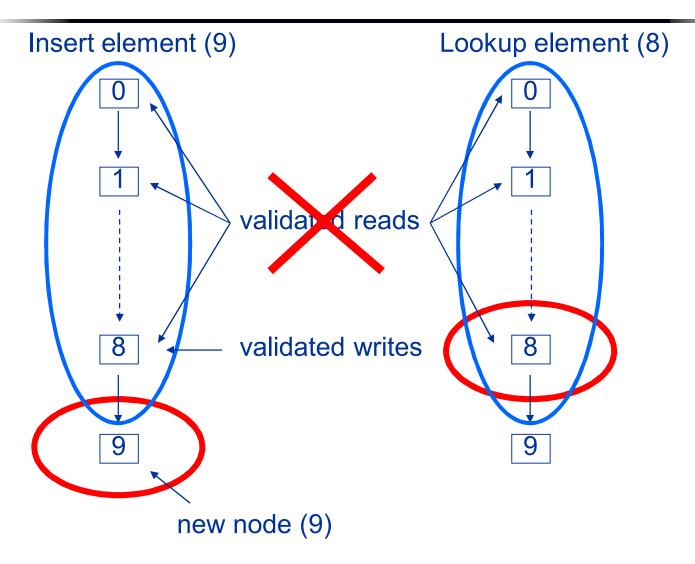
## Validation – Lookup, Insert







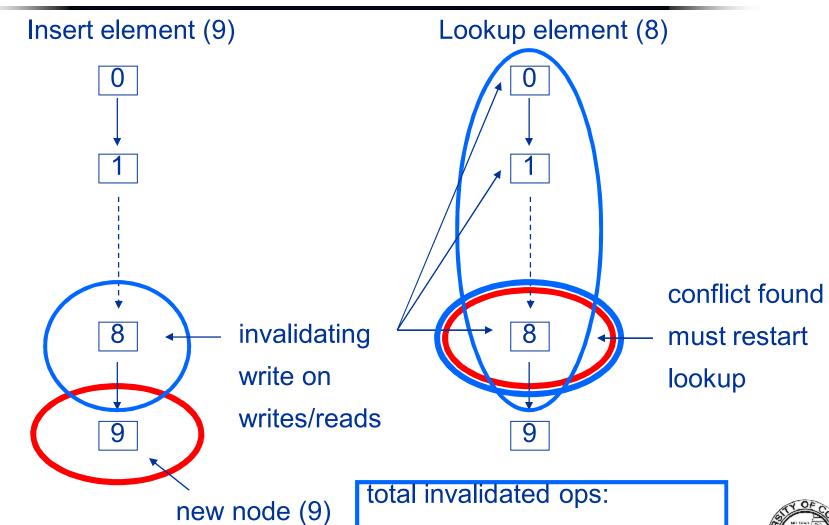
# **Invalidating Linked List**







## Invalidation – Insert, Lookup

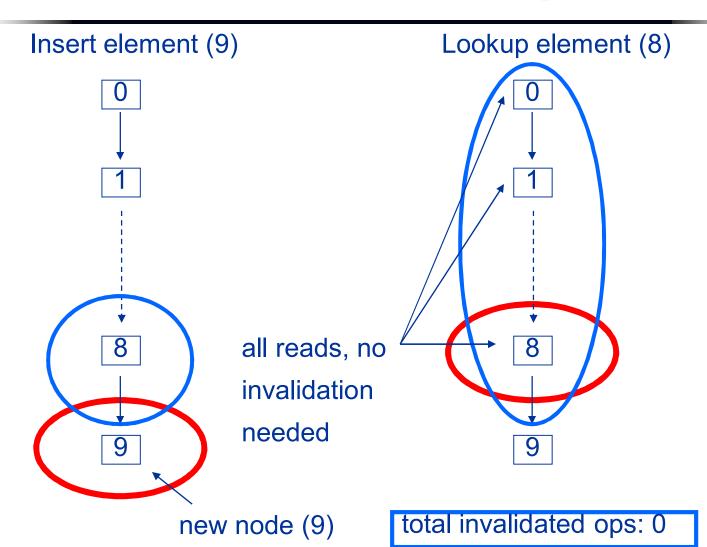




9 reads unopti, 3 reads opti



# Invalidation – Lookup, Insert







# **Consistency Checking**

#### **Validation**

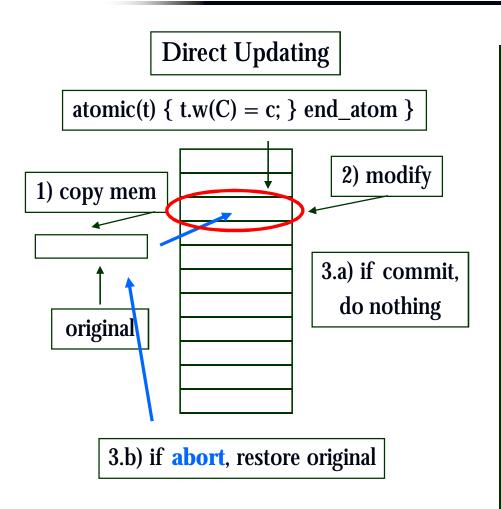
- § Great for short txes with many threads Invalidation
- § Great for long txes with fewer threads
- § Great for read-dominated workloads
  - § Mathematically optimal for read-only txes

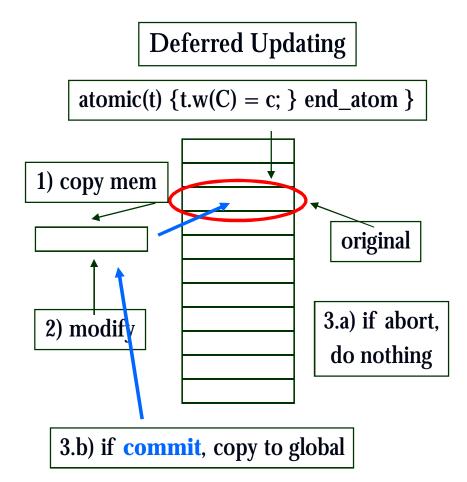
#### TBoost.STM supports both





# **Updating Policies**









# **Updating Policies**

**Direct Updating** 

§ Advantages

fast commit

early writer conflict

detecti

TBoost.STM supports both

§ Disadvantages

single writer per mem speculative contention management **Deferred Updating** 

§ Advantages

fast abort

**truly** optimistic

Letection

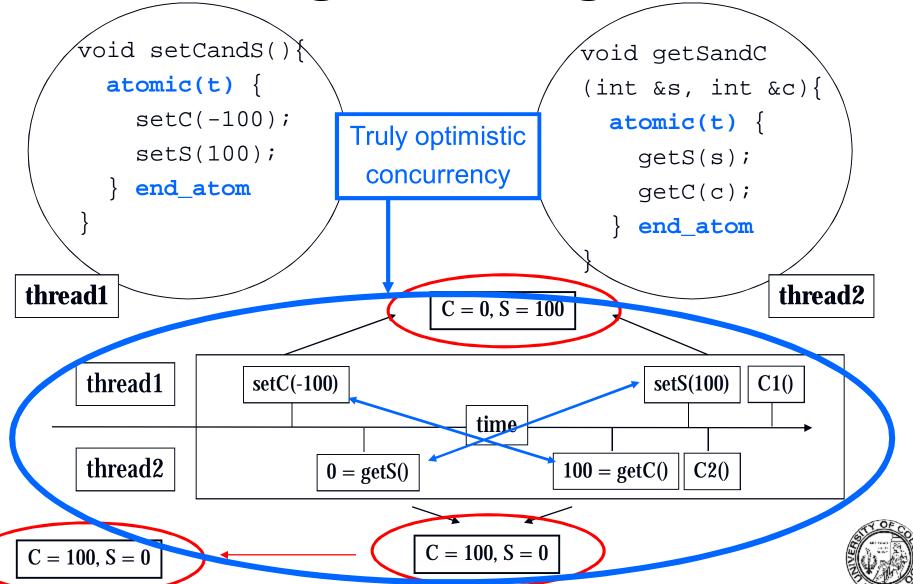
(no speculation)

§ Disadvantages slow commit





Putting It All Together



# Consistency and Updating in TBoost.STM

### Consistency

- § #define PERFORM\_VALIDATION 1
- § #define PERFORM\_VALIDATION 0

## **Updating**

- § transaction::do\_deferred\_updating()
- § transaction::do\_direct\_updating()





## **More Optimizations**

#### TBoost.STM v.0.2

- § Competitive with TL2 (state-of-the-art STM)
- § +10x faster than DracoSTM v.0.1
- § Downloadable from website / boost-sandbox





# Roadmap to boost::stm

## ... And The Big Open Problems







#### TBoost.STM + Boost

#### Code in Boost sandbox

Identify necessary interfaces

- § Future extensions key to success
- § Discussion with Boost community

Goal: review of boost::stm v1.0 early 2010





## TBoost.STM + Research

#### Transactions + locks + others

§ Implementation and formal semantics

## Efficiency

§ Transactional boosting, open nesting, etc.

#### **Test Suites**

- § Integrate with STAMP, STMBench7
  - § STMBench7 researchers helping us now

#### And much more ...





## **Parallel Computing Publications**

- Shifting the Parallel Programming Paradigm (Best Presentation Award) [Raytheon Information Systems and Computing (ISaC), March 2009]
- **Lock-Aware Transactional Memory** [ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), March 2009]
- Optimizing Consistency Checking for Memory-Intensive
  Transactions [ACM Symposium on Principles of Distributed Computing (PODC), August 2008]
- C++ Move Semantics for Exception Safety and Optimization in Software Transactional Libraries [International Workshop on Implementation, Compilation, Optimization of Object-Oriented Languages, Programs and Systems (ICOOOLPS), July 2008]
- Extending Contention Managers for User-Defined Priority-Based
  Transactions [ACM Workshop on Exploiting Parallelism with Transactional
  Memory and other Hardware Assisted Methods (EPHAM), April 2008]
- DracoSTM: A Practical C++ Approach to Software Transactional Memory [ACM Symposium on Library-Centric Software Design (LCSD), October 2007 ]
- Exploration of Lock-Based Software Transactional Memory [MS Thesis, 2007]

  http://eces.colorado.edu/~gottschl

## **Questions?**

# Justin E. Gottschlich (gottschl@colorado.edu)

Department of Electrical and Computer Engineering
University of Colorado-Boulder















